

CS 378: Final Project

Sai Avala
Vaibhav Gupta
Sudheesh Katkam

December 10, 2014

1 Introduction

For this project we wanted to work on a system of linear equations solver. There are various ways of solving a system of linear equations, but for this project we decided to work on the *Jacobi Algorithm*. The way the Jacobi algorithm works is by the choosing a simple solution, and then increasing the number of iterations until you converge to a certain solution. At this point that converged solution should hopefully be the correct answer. There might be a minor error, but this can be controlled by you. The first part of this project was to implement the regular iterative solution to the Jacobi algorithm. We then went ahead and used OpenMP to parallelize it and also vectorized parts of the implementation to further increase the speedup. In addition to implementing a hybrid parallelized and vectorized solution, we also used Valgrind and Gprof in order to profile our code and then improve our solutions. As for reference, our algorithms are optimized for dense matrices.

2 The Algorithm

Here we will describe the basic algorithm. Lets take the standard system of linear equations, $Ax = b$. We are given matrix A , and the vector b , but want to solve for x . The way the iterative Jacobi algorithm works is by first assuming a solution and then iterating until we've converged to an error point. The way we do this is by understanding that matrix A can be rewritten as $A = D + L + U$, where D is the diagonalization matrix and $L + U$ is the summation of the LU factorization. We can then rearrange the formula such that $x_{i+1} = D^{-1}(b - (L + U)x)$.

3 Results

A couple of our results are below:




Figure 1: Varying density of the matrix, outputs runtime

The figure above shows the runtime while varying the density of the matrices. The dotted lines represent matrices that use the vectorized solution and the solid lines are those that don't use the vectorized implementation. As you can see, when we turn on vectorization, we get a significant performance boost. The figure above was run on only 4 thread.




Figure 2: Varying number of threads amongst density

The figure above shows the runtime when varying the number of threads amongst a 70% dense matrix. As you can see, as the number of threads increases, we improve in runtime. We further improve in runtime when we turn on vectorization as well (the dotted lines) in addition to the multithreading.

4 Profiling

While working on this project we used Valgrind and Gprof to help guide us when working on the performance boosts. Valgrind was able to provide information regarding memory usage in the program. We noticed that there was some error in branch prediction, but this is most likely due to using OpenMP. We didn't have an error in branch prediction prior to using OpenMP. We used Gprof to note what function took the most time and was the most expensive for us. By using Gprof, we noticed that generating small matrices took a very long time. However as the matrices increased in size, matrix generation didn't take as long. Our use of `frand()` generally took awhile, because we had to generate random floating number to fill in the matrix. Calculating error took a long time, however we ended up not optimizing this. In addition, we used Gprof to verify that our use of vectorization was actually more optimal than a standard iterative solution. In addition, we used Gprof to minimize general stack usage.

5 Challenges

While working on this project we did run into a couple challenges. In the beginning, we had an issue where our solution wasn't converging in fast enough iterations. What we mean by this is that our solution should have converged at maybe 10 iterations rather than 86 iterations. This was a minor issue that we could solve. The main point of interest is how we handled vectorization and multithreading. We originally used C++11 threads, but to make everything more straightforward in terms of synchronization, we used OpenMP. Our challenge with completing the vectorization was in making sure that computations (for the results) were still the same. Because we were working with heavy floating point numbers, we had to avoid the issue of some small floating point arithmetic error.

6 Future Work

Currently this project uses OpenMP and Intel SSE intrinsics for vectorization. OpenMP is a multithreading framework which makes use of shared memory. However, if we were to take this project further, we would want to try to use MPI. MPI is a message passing interface which allows us to do multicore processing amongst distributed memory. By using MPI, we would want to see if there is even more significant speedups.